



Anna Johnson

Kyle Son

Joe Goldsich

Bill Yerkes

Analysis for determination of a relationship between energy demand and weather.

CSEE5590 – Big Data Programming

Content

- Introduction:
 - Goals and Objectives
 - Motivation
 - Significance
- Iteration 2:
 - Joe Goldsich: Hive - Hadoop - MapReduce (Weather Data Set)
 - Anna Johnson: Sqoop - Hadoop - MySQL
 - Kyle Son: Cassandra - Hive -Hadoop (Merge Table)
 - Bill Yerkes: Hive - Hadoop - MapReduce (Energy Data Set)
- Data Set Link
 - <https://www.kaggle.com/nicholasjhana/energy-consumption-generation-prices-and-weather>

Goals and Objectives

- Utilize the tools and technologies learned from CSEE 5590 to be able to analyze collected data so that it will be possible to determine if there is a relationship between weather and energy consumption and if a relationship exists determine the possibilities of using that relationship to predict future energy needs.

Motivation

The global population continues to increase, and the weather patterns seem to be getting more extreme, from extending periods of both above and below normal temperatures in various parts of the world and in the United States.

The demand and consumption of energy increase with the population and with the extreme weather, the need for air conditioning in the summer and for heating in the winter.

The recent crisis in Texas has demonstrated what can happen if the energy providers are not able to meet the demands of the consumers.

Being able to forecast accurately future demand and plan accordingly can help prevent or mitigate such crises in the future

Significance

- Better planning of resources for Utility Companies can result in reduced cost to the consumers and more reliable service. This also dips into the area of public safety, as loss of power during extreme weather with no warning can be dangerous for vulnerable groups.

Joe Goldsich

- Using Hive, I have practiced making queries on our data sets that might return useful information.
- First, I made simple queries like looking for a correlation between total energy load at different temperature ranges.
- When attempting more complicated queries there was a considerable dip in performance. So I then tried to take advantage of data partitioning that Hive allows.

Grab the energy use at temperatures above 30 C or below 10 C (86F, 50F)

```
SELECT w.city_name, AVG(e.total_load_actual)
FROM weather AS w
INNER JOIN energy AS e
ON w.dt_iso = e.time
WHERE w.temp > 383.15 OR w.temp < 303.15 /*Below 10C higher than 30C (50 - 86*/
GROUP BY w.city_name;
```

This produced:

```
Total MapReduce CPU Time Spent: 8 seconds 920 msec
OK
Barcelona      28316.030604196247
Bilbao 28286.30073713927
Madrid 29125.597647667055
Seville 29662.865883807168
Valencia       28404.889145496534
Time taken: 89.553 seconds, Fetched: 5 row(s)
hive> █
```

And then compare with temperatures between 20 C and 25 C (68 F to 77 F)

```
SELECT w.city_name, AVG(e.total_load_actual)
FROM weather AS w
INNER JOIN energy AS e
ON w.dt_iso = e.time
WHERE w.temp > 393.15 AND w.temp < 398.15 /*Between 20 and 25C (68 - 77)*/
GROUP BY w.city_name;
```

And this produced:

```
J MITC: 129 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 150 ms
OK
Barcelona      28193.09695734459
Bilbao  30800.297554697554
Madrid  28273.834212840808
Seville 28026.180286310766
Valencia      28202.34704
Time taken: 90.513 seconds, Fetched: 5 row(s)
hive> █
```


While attempting to conduct slightly more complex queries (nested selects and joins) ran into some significant performance issues (one query took over 10 minutes to execute). So I then tried to partition the data and see if that made for better performance.

```
hive> CREATE TABLE weather_partition(dt_iso STRING, city_name STRING, temp FLOAT, temp_min F
LOAT, temp_max FLOAT, pressure INT, humidity INT, wind speed INT, wind deg INT, rain_1h FLOA
T, rain_3h FLOAT, snow_3h FLOAT, clouds_all INT, weather_id INT, weather_main STRING, weathe
r_description STRING, weather_icon STRING) PARTITIONED BY(name STRING) row format delimited
fields terminated by ',' stored as textfile;
OK
Time taken: 0.162 seconds
hive> DESC weather_partition;
OK
dt_iso          string
city_name       string
temp            float
temp_min        float
temp_max        float
pressure        int
humidity        int
wind_speed      int
wind_deg        int
rain_1h         float
rain_3h         float
snow_3h         float
clouds_all      int
weather_id      int
weather_main    string
weather_descrip string
weather_icon    string
name            string

# Partition Information
# col_name      data_type      comment
name            string
Time taken: 0.095 seconds, Fetched: 23 row(s)
hive> ALTER TABLE weather_partition ADD PARTITION (name='Valencia');
```

Some partitions examples:

```
hive> ALTER TABLE weather_partition ADD PARTITION (name='July');
```

```
INSERT OVERWRITE TABLE weather_partition PARTITION (name='July') SELECT * FROM weather WHERE month(dt_iso) = 7;
```

< weather_partition

name=April

name=August

name=Barcelona

name=Bilbao

name=Bilboa

name=December

name=February

name=January

name=July

name=June

name=Madrid

name=March

name=May

name=November

name=October

name=September

name=Seville

name=Valencia

View as binary

Download

View file location

Refresh

Last modified 03/22/2021 9:14 PM

User cloudera

Group supergroup

Size 1 51 MB

Home

Page 1 to 50 of 395

user / hive / warehouse / weather_partition / name=September / 000000_0

2015-09-01 00:00:00+02:00, Valencia, 295.76, 292.04, 297.59, 1014, 88, 1, 0, 0.0, 0.0, 0.0, 20, 801, clouds, few clouds, 02n

2015-09-01 01:00:00+02:00, Valencia, 295.47, 292.04, 297.15, 1014, 83, 2, 310, 0.0, 0.0, 0.0, 0, 800, clear, sky is clear, 01n

2015-09-01 02:00:00+02:00, Valencia, 294.64, 290.37, 296.48, 1014, 83, 1, 260, 0.0, 0.0, 0.0, 0, 800, clear, sky is clear, 01n

2015-09-01 03:00:00+02:00, Valencia, 294.47, 290.37, 296.15, 1014, 83, 2, 280, 0.0, 0.0, 0.0, 0, 800, clear, sky is clear, 01n

2015-09-01 04:00:00+02:00, Valencia, 294.23, 290.37, 296.15, 1014, 83, 1, 270, 0.0, 0.0, 0.0, 0, 800, clear, sky is clear, 01n

2015-09-01 05:00:00+02:00, Valencia, 293.94, 289.82, 296.15, 1014, 83, 0, 0, 0.0, 0.0, 0.0, 0, 800, clear, sky is clear, 01n

2015-09-01 06:00:00+02:00, Valencia, 293.46, 289.82, 295.15, 1014, 88, 0, 0, 0.0, 0.0, 0.0, 0, 800, clear, sky is clear, 01n

2015-09-01 07:00:00+02:00, Valencia, 293.02, 289.26, 295.15, 1014, 83, 0, 0, 0.0, 0.0, 0.0, 0, 800, clear, sky is clear, 01n

Anna Johnson

- In MySQL, I created tables for both the energy and weather features datasets
- I populated the tables by loading the data from the csv files
- I used MySQL to run queries on the tables to better understand the data
- I then used Sqoop to import the tables into HDFS
- Sqoop can be used to transfer our databases between the Hadoop/ Hive ecosystem and the relational database system of MySQL, which provides different functionality for querying the data.

Creating Tables in MySQL

```
mysql> create table energy_datas(time VARCHAR(50), biomass FLOAT, lignite FLOAT,
coal_derived_gas FLOAT, fossil_gas FLOAT, hard_coal FLOAT, fossil_oil FLOAT, oi
l_shale FLOAT, peat FLOAT, geothermal FLOAT, hydro_pumped_agg FLOAT, hydro_pumpe
d_consume FLOAT, hydro_run_of_river FLOAT, hydro_water_res FLOAT, marine FLOAT,
nuclear FLOAT, gen_other FLOAT, gen_other_renew FLOAT, solar FLOAT, waste FLOAT,
gen_wind_offshore FLOAT, gen_wind_onshore FLOAT, forecast_solar FLOAT, forecast
_wind_offshore FLOAT, forecast_wind_onshore FLOAT, total_load_forecast FLOAT, to
tal_load_actual FLOAT, price_day_ahead FLOAT, price_actual FLOAT);
Query OK, 0 rows affected (0.02 sec)
```

```
mysql> create table weather_data(dt_iso VARCHAR(100), city_name VARCHAR(100), te
mp FLOAT, temp_min FLOAT, temp_max FLOAT, pressure INT, humidity INT, wind speed
INT, wind_deg INT, rain_1h FLOAT, rain_3h FLOAT, snow_3h FLOAT, clouds_all INT,
weather_id INT, weather_main VARCHAR(100), description VARCHAR(100), weather_ic
on VARCHAR(20));
Query OK, 0 rows affected (0.04 sec)
```

Populating tables with data:

```
mysql> load data local infile '/home/cloudera/Desktop/weather_features.csv' into
table weather_test
-> fields terminated by ','
-> lines terminated by '\n';
Query OK, 178397 rows affected, 12 warnings (2.94 sec)
Records: 178397 Deleted: 0 Skipped: 0 Warnings: 6
```

Example Queries in MySQL:

Weather data:

```
mysql> SELECT AVG(temp), MAX(temp), MIN(temp), city_name FROM weather_data GROUP BY city_name;
```

AVG(temp)	MAX(temp)	MIN(temp)	city_name
289.848244622644	309.15	262.24	Barcelona
286.378488296441	312.47	266.85	Bilbao
0	0	0	city_name
288.061070234593	313.33	264.132	Madrid
293.105430544879	315.6	271.05	Seville
290.780776819068	311.15	268.831	Valencia

```
6 rows in set (0.94 sec)
```

Energy data:

```
mysql> SELECT AVG(waste), MAX(waste), MIN(waste) FROM energy_datas;
```

AVG(waste)	MAX(waste)	MIN(waste)
269.298445743619	357	0

1 row in set (0.04 sec)

```
mysql> SELECT time, waste FROM energy_datas WHERE waste > 269 OR waste = 0 limit 10;
```

time	waste
time	0
2015-01-05 03:00:00+01:00	0
2015-01-05 12:00:00+01:00	0
2015-01-05 13:00:00+01:00	0
2015-01-05 14:00:00+01:00	0
2015-01-05 15:00:00+01:00	0
2015-01-05 16:00:00+01:00	0
2015-01-05 17:00:00+01:00	0
2015-01-09 00:00:00+01:00	273
2015-01-09 19:00:00+01:00	281

10 rows in set (0.00 sec)

```
mysql> SELECT time, solar, fossil_oil FROM energy_datas WHERE solar > fossil_oil limit 10;
```

time	solar	fossil_oil
2015-01-01 09:00:00+01:00	743	163
2015-01-01 10:00:00+01:00	2019	167
2015-01-01 11:00:00+01:00	3197	166
2015-01-01 12:00:00+01:00	3885	167
2015-01-01 13:00:00+01:00	4007	167
2015-01-01 14:00:00+01:00	3973	166
2015-01-01 15:00:00+01:00	3818	160
2015-01-01 16:00:00+01:00	3088	163
2015-01-01 17:00:00+01:00	1467	165
2015-01-01 18:00:00+01:00	404	164

10 rows in set (0.00 sec)

Transferring Data with Sqoop

Importing from MySQL to HDFS:

```
[cloudera@quickstart ~]$ sqoop import --connect jdbc:mysql://localhost/weather -  
-username root --password cloudera --table weather_data --m 1
```

Exporting between Hive and MySQL:

```
[cloudera@quickstart ~]$ sqoop export --connect jdbc:mysql://localhost/weather -  
-username root --password cloudera --table weather_test --export-dir /user/hive/  
warehouse/weather_dbh.db/weather_datah --input-fields-terminated-by ',' --input-  
lines-terminated-by '\n' -m 1
```

energy_datas and weather_data in HDFS after importing from MySQL:

21/03/22 11:40:45 INFO mapreduce.ImportJobBase: Retrieved 35065 records.

```
[cloudera@quickstart ~]$ hadoop fs -ls
```

Found 7 items

drwxr-xr-x	-	cloudera	cloudera	0	2021-02-25	19:28	acad
drwxr-xr-x	-	cloudera	cloudera	0	2021-03-22	11:40	energy_datas
drwxr-xr-x	-	cloudera	cloudera	0	2021-01-27	18:42	johnsona9726
drwxr-xr-x	-	cloudera	cloudera	0	2021-02-25	17:41	queries
drwxr-xr-x	-	cloudera	cloudera	0	2021-02-25	16:46	queryresult
drwxr-xr-x	-	cloudera	cloudera	0	2021-03-21	21:49	weather_data
.

Kyle Son

- Instead of Hive CLI, I used HIVE beeline for better visualization of table
- In Hive, Join two tables on (dt_iso = time) and perform queries on it
- Using cassandra, perform some queries on each table
- It is not possible to join two tables in cassandra
- I used sparksql on intellij with scala for query the data

HIVE(Beeline)

```
hive> CREATE TABLE Energy (time TIMESTAMP,generation_biomass FLOAT,generation_fossil_brown_coal_lignite
FLOAT,generation_fossil_coal_derived_gas FLOAT,generation_fossil_gas FLOAT,generation_fossil_hard_coal F
LOAT,generation_fossil_oil FLOAT,generation_fossil_oil_shale FLOAT,generation_fossil_peat FLOAT,generati
on_geothermal FLOAT,generation_hydro_pumped_storage_aggregated FLOAT,generation_hydro_pumped_storage_con
sumption FLOAT,generation_hydro_run_of_river_and_poundage FLOAT,generation_hydro_water_reservoir FLOAT,g
eneration_marine FLOAT,generation_nuclear FLOAT,generation_other FLOAT,generation_other_renewable FLOAT,
generation_solar FLOAT,generation_waste FLOAT,generation_wind_offshore FLOAT,generation_wind_onshore FLO
AT,forecast_solar_day_ahead FLOAT,forecast_wind_offshore_eday_ahead FLOAT,forecast_wind_onshore_day_ahea
d FLOAT,total_load_forecast FLOAT,total_load_actual FLOAT,price_day_ahead FLOAT,price_actual FLOAT
> )
> row format delimited fields terminated by ','
> stored AS textfile
> tblproperties("skip.header.line.count"="1");
OK
Time taken: 0.236 seconds
hive> CREATE TABLE Weather (dt_iso TIMESTAMP,city_name STRING,temp FLOAT,temp_min FLOAT,temp_max FLOAT,p
ressure INT,humidity INT,wind_speed INT,wind_deg INT,rain_1h FLOAT,rain_3h FLOAT,snow_3h FLOAT,clouds_al
l FLOAT,weather_id INT,weather_main STRING,weather_description STRING,weather_icon STRING
> )
> row format delimited fields terminated by ','
> stored AS textfile
> tblproperties("skip.header.line.count"="1");
OK
Time taken: 0.089 seconds
```

```
[cloudera@quickstart hive-1.1.0-cdh5.13.0+1269]$ beeline;
Beeline version 1.1.0-cdh5.13.0 by Apache Hive
beeline> show databases;
No current connection
beeline> !connect jdbc:hive2://
scan complete in 4ms
Connecting to jdbc:hive2://
Enter username for jdbc:hive2://: cloudera
Enter password for jdbc:hive2://: *****
Connected to: Apache Hive (version 1.1.0-cdh5.13.0)
Driver: Hive JDBC (version 1.1.0-cdh5.13.0)
Transaction isolation: TRANSACTION_REPEATABLE_READ
0: jdbc:hive2://> show databases;
OK
+-----+-----+
| database_name |
+-----+-----+
| bigdataproject |
| db1 |
| default |
+-----+-----+
3 rows selected (1.806 seconds)
0: jdbc:hive2://>
```

create table merged as select * from (select * from Energy e left join Weather w on e.time = w.dt_iso)x

=> From This query, I created a merged table of energy and weather

```
SELECT city_name, AVG(temp) as AvgTemp, AVG(wind_speed) as AvgWindSpeed, AVG(humidity) as AvgHumidity
from merged group by city_name;
```

```
total number of time spent: 3 seconds 000 msec
OK
+-----+-----+-----+-----+
| city_name | avgtemp | avgwindspeed | avghumidity |
+-----+-----+-----+-----+
| Barcelona | 289.8482446226438 | 2.786588115909347 | 73.99422144548427 |
| Bilbao | 286.3784882964413 | 1.9574698895719174 | 79.08945509165252 |
| Madrid | 288.0610702345934 | 2.4416963079383462 | 59.776932197314366 |
| Seville | 293.1054305448794 | 2.4837865961695305 | 64.14073178277133 |
| Valencia | 290.7807768190682 | 2.6928154787309717 | 65.14511310285958 |
+-----+-----+-----+-----+
5 rows selected (38.935 seconds)
```

=> From This query, I displayed each city's average temperature windspeed and humidity group by city name

```
SELECT x.city_name AS city_name, x.year AS year, AVG(x.price_actual) AS avg_price_actual,
AVG(x.price_day_ahead) AS avg_price_ahead
FROM (SELECT city_name, YEAR(time) AS YEAR, price_actual, price_day_ahead FROM merged) AS x
GROUP BY city_name, year;
```

OK

city_name	year	avg_price_actual	avg_price_ahead
Barcelona	2015	61.37339986352105	50.34858796293164
Barcelona	2016	47.408718672313334	39.706185255424245
Barcelona	2017	59.336550360126544	52.264069580019545
Barcelona	2018	63.41520793983889	57.24723606680608
Bilbao	2015	61.43668094638449	50.40329749004072
Bilbao	2016	47.504653980321336	39.73472841074056
Bilbao	2017	59.43972139166336	52.22414713690373
Bilbao	2018	63.37481516512984	57.217914237181084
Madrid	2015	61.3486871444867	50.32945377248953
Madrid	2016	47.70215140951364	39.848420415028215
Madrid	2017	59.489210907216496	52.3455380426181
Madrid	2018	63.475907996585576	57.38593137540892
Seville	2015	61.34763784962553	50.3336542601593
Seville	2016	47.44788847544184	39.619593737169
Seville	2017	59.30196663987525	52.225143714592235
Seville	2018	63.368107600977105	57.221872207268525
Valencia	2015	61.36173105451796	50.33310753950027
Valencia	2016	47.43544185960827	39.68283255650655
Valencia	2017	59.32627164795927	52.24235188971896
Valencia	2018	63.45932222534277	57.31257539540139

20 rows selected (39.679 seconds)

=> From This query, I displayed avg actual price and ahead price group by city and year to show the chronological change of the electricity price

CASSANDRA

```
jk@jk-VirtualBox:~$ cqlsh
Connected to Test Cluster at 127.0.0.1:9042.
[cqlsh 5.0.1 | Cassandra 3.9 | CQL spec 3.4.2 | Native protocol v4]
Use HELP for help.
cqlsh> describe keyspaces;

system_schema  system              system_traces  data
system_auth    system_distributed  bigdataproject

cqlsh> use bigdataproject;
cqlsh:bigdataproject> describe tables;

energy weather

cqlsh:bigdataproject>
```

```
cqlsh:bigdataproject> select * from weather where temp >290 AND city_name = 'Seville' ALLOW FILTERING;
```

dt_iso	city_name	clouds_all	humidity	pressure	rain_1h	rain_3h	snow_3h	temp	temp_max	temp_min	weather_desc
2018-05-31 12:00:00.000000+0000	Seville	0	68	1019	0	0	0	290.54001	291.14999	290.14999	sky
clear 01d 800	clear	30	2								
2018-07-07 17:00:00.000000+0000	Seville	0	26	1015	0	0	0	306.54001	307.14999	306.14999	sky
clear 01d 800	clear	310	2								
2017-05-30 13:00:00.000000+0000	Seville	20	73	1016	0	0	0	294.54001	295.14999	294.14999	fe
clouds 02d 801	clouds	20	0								
2016-11-15 10:00:00.000000+0000	Seville	0	81	1024	0	0	0	290.06	300.14999	282.14999	sky
clear 01d 800	clear	70	4								
2016-06-12 07:00:00.000000+0000	Seville	8	70	1023	0	0	0	290.289	290.289	290.289	sky
clear 02n 800	clear	7	1								
2017-04-23 15:00:00.000000+0000	Seville	0	43	1016	0	0	0	300.22	309.14999	294.14999	sky
clear 01d 800	clear	156	1								
2018-08-27 00:00:00.000000+0000	Seville	0	33	1011	0	0	0	305.14999	305.14999	305.14999	sky
clear 01n 800	clear	230	3								
2016-03-05 18:00:00.000000+0000	Seville	0	36	1014	0	0	0	292.44	298.14999	289.14999	sky
clear 01d 800	clear	320	5								
2018-10-02 07:00:00.000000+0000	Seville	40	83	1017	0	0	0	292.94	294.14999	292.14999	scattered
clouds 03n 802	clouds	30	2								
2017-08-30 13:00:00.000000+0000	Seville	20	77	1016	0	0	0	293.54001	294.14999	293.14999	fe
clouds 02d 801	clouds	0	0								
2015-10-05 09:00:00.000000+0000	Seville	20	69	1019	0	0	0	298.25	300.14999	296.14999	
mist 50n 701	mist	200	3								
2016-07-06 06:00:00.000000+0000	Seville	0	78	1017	0	0	0	298.82999	306.14999	294.14999	sky
clear 01n 800	clear	210	3								
2017-03-06 16:00:00.000000+0000	Seville	0	56	1027	0	0	0	294.64999	299.14999	291.14999	sky
clear 01d 800	clear	0	0								

=> Created two separate table in cassandra, and perform query for each tables

```
cqlsh:bigdataproyect> select time, total_load_forecast, total_load_actual, price_day_ahead, price_actual from  
m energy where price_actual >70 LIMIT 10 ALLOW FILTERING;
```

time	total_load_forecast	total_load_actual	price_day_ahead	price_actual
2015-01-08 19:00:00.000000+0000	28124	27507	52.04	90.87
2015-08-24 12:00:00.000000+0000	31216	30648	62.47	70.38
2015-01-06 21:00:00.000000+0000	29307	30120	57	81.65
2015-02-12 16:00:00.000000+0000	32620	32607	66.4	80.05
2015-03-06 11:00:00.000000+0000	32292	32003	60.21	70.29
2015-07-21 01:00:00.000000+0000	32205	32737	63	75.78
2015-02-05 19:00:00.000000+0000	23765	24216	32.08	75.71
2015-12-23 10:00:00.000000+0000	33123	32680	64.54	71.72
2015-08-03 12:00:00.000000+0000	25938	25789	38.17	77.98
2015-06-27 15:00:00.000000+0000	31866	31970	65	74.87

=> From This query, I only display the data where actual price is over 70

SPARKSQL(IntelliJ)

```
GroupProject.scala x build.sbt x

def main(args: Array[String]): Unit = {

    val spark: SparkSession = SparkSession.builder()
        .master("local[*]")
        .appName("groupProject")
        .getOrCreate()

    val filepath1 = "energy_dataset.csv"
    val filepath2 = "weather_features.csv"
    val df_energy = spark.read.options(Map("inferSchema" -> "true", "sep" -> ",", "header" -> "true")).csv(filepath1)
    val df_weather = spark.read.options(Map("inferSchema" -> "true", "sep" -> ",", "header" -> "true")).csv(filepath2)

    val df_merge = df_energy.join(df_weather, df_energy("time") === df_weather("dt_iso"), joinType = "inner")
    df_merge.show()

    val df_weatherby = df_merge.groupBy(col => "weather_description")
        .avg(colNames = "price actual", "price day ahead", "total load actual")
    df_weatherby.show()
}
```

weather_description	avg(price actual)	avg(price day ahead)	avg(total load actual)
fog	61.98602154828406	51.95822426177173	27938.253391859536
drizzle	56.84691056910567	51.23531165311652	29612.173441734416
very heavy rain	55.3823076923077	46.016666666666666	27669.25641025641
ragged shower rain	68.61	50.6	26513.0
proximity shower ...	56.882668067226895	52.44361344537816	30394.945263157893
light thunderstorm	58.455	52.36	28351.5
few clouds	57.73782467835898	50.286180181303074	29256.7720938932
heavy intensity s...	57.78037037037037	53.74283950617283	29858.0987654321
proximity moderat...	62.74	56.347500000000004	29026.0
haze	51.93441379310345	42.21977011494253	25578.075862068967
shower sleet	11.65	4.0	25136.0
light rain	55.99469234296196	48.28205226960125	28280.440106558883
dust	58.34052173913043	49.90831884057972	29306.6
light intensity d...	56.87957292506042	51.253650282030605	29254.7751813054
light intensity s...	58.42899543378993	52.71281582952817	29459.219178082192
proximity thunder...	62.33487500000001	55.290854166666669	28945.922916666666
broken clouds	56.80642824392482	49.646100985786504	28832.29800412939
overcast clouds	57.0338110113237	48.5816516985552	28101.91484375
squalls	70.53	62.16	35513.0
proximity drizzle	64.9	57.41	31462.0

=> From This query, I group by weather description and display price information to show the correlation between them

Bill Yerkes

- Utilized Hive (Hadoop and MapReduce) to analyse the Energy Data
 - Created Tables
 - Imported Data
 - Ran Queries (MapReduce) to get metrics
- Technologies not utilized/required (May used them in the future)
 - Cassandra (NoSQL DB: do not foresee using.)
 - Sqoop (Data Migration: probably will use to move to SQL DB for reports.)
 - Solr (Search Engine: do not foresee using.)

Bill Yerkes

```
CREATE TABLE Weather (  
dt iso TIMESTAMP,  
city_name STRING,  
temp DOUBLE,  
temp_min DOUBLE,  
temp_max DOUBLE,  
pressure INT,  
humidity INT,  
wind_speed INT,  
wind_deg INT,  
rain_1h DOUBLE,  
rain_3h DOUBLE,  
snow_3h DOUBLE,  
clouds_all INT,  
weather_id INT,  
weather_main STRING,  
weather_description STRING,  
weather_icon STRING)  
row format delimited fields terminated by ',' stored as textfile;  
  
load data local inpath '/home/cloudera/Downloads/weather_features.csv' into table Weather;  
  
load data local inpath '/home/cloudera/Downloads/energy_dataset.csv' into table Energy;
```

```
CREATE TABLE Energy (  
time TIMESTAMP,  
generation_biomass INT,  
generation_fossil INT,  
brown_coal_lignite INT,  
generation_fossil_coal_derived_gas INT,  
generation_fossil_gas INT,  
generation_fossil_hard_coal INT,  
generation_fossil_oil INT,  
generation_fossil_oil_shale INT,  
generation_fossil_peat INT,  
generation_geothermal INT,  
generation_hydro_pumped_storage_aggregated INT,  
generation_hydro_pumped_storage_consumption INT,  
generation_hydro_run_of_river_and_poundage INT,  
generation_hydro_water_reservoir INT,  
generation_marine INT,  
generation_nuclear INT,  
generation_other INT,  
generation_other_renewable INT,  
generation_solar INT,  
generation_waste INT,  
generation_wind_offshore INT,  
generation_wind_onshore INT,  
forecast_solar_day_ahead INT,  
forecast_wind_offshore_day_ahead INT,  
forecast_wind_onshore_day_ahead INT,  
total_load_forecast INT,  
total_load_actual INT,  
price_day_ahead DOUBLE,  
price_actual DOUBLE)  
row format delimited fields terminated by ',' stored as textfile;
```


Bill Yerkes

Generate Metrics via Hive (MapReduce)

Select AVG(generation_biomass),MAX(generation_biomass),MIN(generation_biomass), STDDEV(generation_biomass) FROM energy;

```
[hive> Select AVG(generation_biomass),MAX(generation_biomass), STDDEV(generation_biomass) FROM energy;
Query ID = cloudera_20210322081212_cf72c18b-fcc7-4ba9-90b5-9b1fad603eeb
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1614704973316_0010, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1614704973316_0010/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1614704973316_0010
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-03-22 08:12:54,412 Stage-1 map = 0%, reduce = 0%
2021-03-22 08:13:08,748 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.24 sec
2021-03-22 08:13:25,130 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.98 sec
MapReduce Total cumulative CPU time: 3 seconds 980 msec
Ended Job = job_1614704973316_0010
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.98 sec HDFS Read: 6288615 HDFS Write: 41 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 980 msec
OK
383.51353973462693      592      85.35272526880253
Time taken: 47.154 seconds, Fetched: 1 row(s)
hive>
```

	Average	Max	Min	STDV
generation biomass	383.51	592.00	0.00	85.35
generation fossil brown coal/lignite	448.06	999.00	0.00	354.57
generation fossil gas	5622.74	20034.00	0.00	2201.83
generation fossil hard coal	4256.07	8359.00	0.00	1961.60
generation fossil oil	298.32	449.00	0.00	52.52
generation hydro pumped storage consumption	475.58	4523.00	0.00	792.41
generation hydro run-of-river and poundage	972.12	2000.00	0.00	400.78
generation hydro water reservoir	2605.11	9728.00	0.00	1835.20
generation nuclear	6263.91	7117.00	0.00	839.67
generation other	60.23	106.00	0.00	20.24
generation other renewable	85.64	119.00	0.00	14.08
generation solar	1432.67	5792.00	0.00	1680.12
generation waste	269.45	357.00	0.00	50.20
generation wind onshore	5464.48	17436.00	0.00	3213.69
forecast wind onshore day ahead	5471.22	17430.00	237.00	3176.31
total load forecast	28712.13	41390.00	18105.00	4594.10
total load actual	28696.94	41015.00	18041.00	4574.99
price day ahead	49.87	101.99	2.06	14.62
price actual	57.88	116.80	9.33	14.20

Next Steps

- Group will pool information about analysis of Data.
- Investigate migration of current functions from Hive to Spark.
- Investigate future class technologies on how to perform predictions and generate visualizations of data.
- Join datasets based on time column and look for the relationships between the two

Q & A

THANK
YOU